

iTrace: an algorithmically verified open-source core for webcam gaze, saccade, and pupil analysis

Daniel Ari Friedman (Active Inference Institute, FractAI)

2026-05-29

1 Abstract

Webcam-based eye tracking promises low-cost access to gaze, saccade, and pupil signals, but the open-source landscape is fragmented across browser gaze trackers, appearance models, event classifiers, pupillometry packages, and capture shells whose hardware dependencies make reproducible, headless validation difficult. We present **iTrace**, a Python toolkit that separates a pure NumPy/SciPy analysis core — gaze geometry, Savitzky–Golay velocity, I-VT and I-DT event identification (Salvucci and Goldberg 2000), Engbert–Kliegl microsaccade detection (Engbert and Kliegl 2003), saturating and power-law main-sequence fitting (Bahill et al. 1975), causal real-time pupil-phase detection (Kronemer et al. 2024), and scanpath direction encoding — from an optional, thin webcam/MediaPipe capture shell. Because the core never imports a hardware dependency, the entire test suite runs on a headless machine and algorithmically verifies each detector against synthetic signals whose ground truth is known by construction: the I-VT detector recovers a 10–12° saccade’s amplitude to within 5% and its peak velocity to within 10%, and a second, independently-formulated implementation cross-checks the event boundaries. A 3-D eyeball forward model closes the loop — projecting a known gaze through a pinhole camera to landmarks and back through the estimator — recovering gaze to an RMS residual of 0.16°, an internal-consistency check that exercises the geometry path (not just the event algorithms) and bounds, but does not establish, real-world error. A Monte-Carlo noise-sensitivity sweep then varies an idealized landmark-localisation noise σ and quantifies how each signal degrades: gaze stays usable to $\sigma \approx 0.005$ (≈ 3 px at 640 px width) while saccade detection — resting on the noise-amplifying velocity — collapses first at $\sigma \approx 0.0014$ (≈ 0.9 px), implying that webcam saccade timing is theoretically marginal before real MediaPipe, lens, illumination, head-pose, and calibration errors are added. (The pupil channel’s robustness is set by a modelling assumption, not measured, and is reported as conditional.) The package ships with a typer command-line interface, a Streamlit dashboard, a local HTML/WebSocket live orchestrator, and publication-figure scripts, all under a green test suite with 429 tests at 91.18% coverage. iTrace demonstrates that webcam eye-movement software becomes more auditable when the verifiable algorithms are decoupled from fragile hardware.

2 Introduction

Eye movements are a rich, non-invasive window onto attention, cognition, and oculomotor health. Three signal families dominate research use: the **gaze trajectory** (where the eyes point over time), **saccades** (the fast ballistic movements between fixations, with characteristic direction, amplitude, and peak-velocity dynamics), and **pupil diameter** (a marker of arousal and cognitive load). Dedicated infrared eye-trackers measure all three at high precision, but their cost restricts access. A maturing open-source ecosystem now estimates related signals from commodity cameras: browser trackers such as WebGazer (Papoutsaki et al. 2016), large webcam and mobile gaze datasets such as GazeCapture (Krafka et al. 2016) and MPIIGaze (Zhang et al. 2017), appearance-based gaze estimators such as L2CS-Net (Abdelrahman et al. 2022), landmark-based capture stacks such as MediaPipe Iris (Vakunov et al. 2020), velocity- and dispersion-threshold event classifiers (Salvucci and Goldberg 2000; Nyström and Holmqvist 2010; Dar et al. 2021), microsaccade detectors (Engbert and Kliegl 2003), file-oriented processing and quality-reporting pipelines (Krakowczyk et al. 2023; Jakobi et al. 2024), and webcam or offline pupillometry tools (Shah et al. 2025; Zhang and Jonides 2026; Mittner 2020; Cai et al. 2024).

Two problems recur across these tools. First, the *capture* layer (camera drivers, MediaPipe, browser WebRTC (Papoutsaki et al. 2016), integrated open-source platforms such as Pupil (Kassner et al. 2014)) is tightly coupled to the *analysis* layer, so the scientific algorithms cannot be exercised or checked without hardware — which makes continuous integration, regression testing, and reproducible benchmarking awkward. Second, “validation” is often demonstrated on recorded data without a ground-truth oracle, so a detector that silently mis-handles noise, blinks, or NaNs can pass unnoticed.

The existing tools each occupy a different point in this design space, and iTrace is positioned deliberately against them. WebGazer (Papoutsaki et al. 2016) pioneered in-browser webcam gaze but fuses capture, calibration, and estimation into one runtime. GazeCapture and MPIIGaze provide the public-data substrate for real-frame gaze evaluation (Krafka et al.

2016; Zhang et al. 2017), while L2CS-Net demonstrates how appearance models can predict pitch/yaw from unconstrained images (Abdelrahman et al. 2022). pymovements (Krakowczyk et al. 2023) offers a mature, file-oriented processing pipeline and now anchors explicit data-quality reporting standards (Jakobi et al. 2024), but it assumes the gaze samples already exist. REMoDNaV (Dar et al. 2021) is a strong event classifier whose four-class fixation / saccade / smooth-pursuit / post-saccadic-oscillation output is richer than ours, but it classifies an input signal rather than reasoning about how that signal was estimated from a camera. PupilSense/EyeDentify (Shah et al. 2025) addresses webcam pupil diameter with learned models and a dataset; PupEyes and pypillometry focus on reproducible pupil preprocessing and analysis after samples have been recorded (Zhang and Jonides 2026; Mittner 2020); Open-DPSM models luminance-driven pupil responses in dynamic stimuli (Cai et al. 2024). iTrace does not try to out-measure any of these on real eyes. It occupies a narrower gap: a *verification-first*, hardware-decoupled implementation of canonical gaze, saccade, microsaccade, and pupil algorithms whose correctness is pinned to constructed ground truth and whose outputs can later be compared against these datasets and packages.

iTrace addresses both problems. We implement the canonical algorithms in a pure NumPy/SciPy core with no hardware dependency, and we algorithmically verify every detector against *synthetic signals whose events are known by construction*. The capture of webcam frames and MediaPipe iris landmarks is an optional, thin shell whose landmark-to-gaze mathematics is itself testable without a camera. Layered on the verified core is an analysis surface — descriptive event statistics, maximum-likelihood distribution fitting and model comparison, scanpath spread and entropy metrics, a bootstrap confidence interval on the main-sequence exponent, and a deterministic publication-figure gallery — that operates on whatever gaze it is given, synthetic or recorded, without ever reaching for the camera. The result is a toolkit whose scientific methods are auditable precisely because they are decoupled from the fragile hardware. sec. 3 describes the architecture and algorithms; sec. 13 reports the ground-truth verification; sec. 20 states the accuracy limits that constrain webcam eye tracking.

3 Methods: architecture

iTrace is organised as two strata with a one-way dependency: a pure analysis **core** and an optional hardware **shell**. The split is not cosmetic — it is the mechanism that makes the scientific claims of this paper checkable.

The **core** depends only on NumPy, SciPy, and pandas. It divides into a signal-and-event layer (**geometry**, **calibration**, **velocity**, **saccades**, **mainsequence**, **pupil**, **pupilphase**, **encoding**, **eyemodel**, **scene**, **power**, **pipeline**, **io**) and a quantitative **stats** subpackage layered on top of the detected events (**stats.descriptive**, **stats.distributions**, **stats.scanpath_metrics**, **stats.similarity**, **stats.timeseries**; sec. 11, sec. 12). Every value the core consumes or produces is a plain array or a typed dataclass; it has no knowledge of cameras, files beyond CSV, or rendering. The unit and coordinate conventions are fixed at the boundary — pixels with a top-left origin, degrees of visual angle with a mathematical-convention direction (0deg right, +90deg up), seconds for time, and a pupil unit that travels with every pupil number — so no ambiguous quantity ever leaks between layers.

Calibration is kept in that same core layer because it is an analysis transform, not a camera driver. The shipped model is deliberately inspectable — a fitted two-dimensional affine map from raw gaze coordinates to known target coordinates, with RMS, percentile, and maximum error summaries. That makes screen-space calibration auditable when target points exist, while leaving live webcam output honestly labelled as relative/algorithmic when no external target or reference device is present.

The **shell** (**capture**, **dashboard**, **cli**) is the only code that touches OpenCV, MediaPipe, Streamlit, Plotly, or matplotlib, and it imports those dependencies **lazily**, **inside functions**, never at module load; a missing optional dependency raises an actionable error rather than breaking **import itrace**. Two consequences follow, and they are the load-bearing design decisions of the whole package:

1. **import itrace** and the *entire* test suite succeed on a headless machine with none of the optional dependencies installed.
2. The scientific algorithms can therefore be verified continuously, in CI, without a webcam — which is precisely what makes their correctness auditable (sec. 17).

The **stats**, **viz**, and **method** layers obey a second, finer rule that keeps the headless guarantee honest: matplotlib is treated like any other optional plotting backend and lives **only** in the visualisation modules, imported at the top of those modules with the **Agg** backend selected before **pyplot**. A method or statistics module never imports a plotting library, so a fitted distribution, a scanpath-similarity matrix, or a windowed event-rate series is computed and returned as arrays regardless of whether any display is available.

The dependency rule is strict and one-directional: the shell may import the core, the core may never import the shell. A new capture backend (browser WebRTC, a deep-learning gaze model, a research IR camera) is added by writing a shell (sec. 8) that emits the same **GazeStream** / **PupilStream** the core already analyses; the verified core never changes, and every analysis in sec. 5 through sec. 12 applies unaltered to its output. The full module map and dependency contract are in sec. 25.

4 Methods: gaze geometry

The geometry layer turns raw image measurements into calibrated quantities in degrees of visual angle (dva). Pixel offsets from screen centre are converted by the exact relation

$$\theta = \arctan\left(\frac{s_{\text{cm}}}{d_{\text{cm}}}\right) \quad (1)$$

where s_{cm} is the offset in centimetres (pixels times the screen’s cm-per-pixel) and d_{cm} the viewing distance. The inverse of eq. 1, `deg2pix`, is exact, so `pix2deg``deg2pix` round-trips to floating-point tolerance.

For appearance-based capture, normalised iris displacement within the eye aperture is mapped to eyeball rotation through a sphere model,

$$\theta_{\text{gaze}} = \arcsin(o \cdot \sin \theta_{\text{max}}) \quad (2)$$

with $o \in [-1, 1]$ the normalised iris offset and θ_{max} the angle at full deflection; the map (eq. 2) is monotone and odd. Head-distance dependence is removed by dividing offsets by the inter-ocular distance in pixels and rescaling to a fixed reference, so the same physical movement yields the same value regardless of how far the subject sits from the camera.

Two conventions are fixed package-wide to prevent sign errors: gaze direction uses $0^\circ = \text{right}$ and $+90^\circ = \text{up}$ (image- y , which grows downward, is negated on the way in); all non-finite inputs raise rather than silently coercing to zero, so a NaN can never masquerade as a valid centred gaze.

5 Methods: velocity and event detection

5.1 Velocity

Gaze velocity is estimated two ways, matching how data arrives. For uniformly sampled streams a Savitzky–Golay filter fits a local polynomial and returns its analytic first derivative, smoothing and differentiating in one pass; the window is clamped to an odd length not exceeding the trace. For non-uniform timestamps a central-difference gradient is taken on the actual time vector. Two-dimensional speed is the Euclidean norm of the per-axis velocities.

5.2 Fixations and saccades (I-VT, I-DT)

I-VT (Salvucci and Goldberg 2000) labels each sample saccadic when its 2-D speed exceeds a velocity threshold and collapses contiguous same-label runs into events; runs shorter than a minimum duration are reabsorbed into the surrounding fixation, suppressing single-sample velocity spikes. In noisy or intermittently tracked streams the configurable `merge_gap_s` parameter optionally bridges short subthreshold holes inside one high-velocity movement before the duration filter is applied; this prevents a brief landmark dropout or low-velocity plateau from fragmenting one saccade into several events while leaving the historical default at zero for exact backward compatibility. **I-DT** (Salvucci and Goldberg 2000) greedily extends a window while its spatial dispersion $(\max x - \min x) + (\max y - \min y)$ stays below threshold, emitting one fixation per maximal low-dispersion window.

5.3 Microsaccades (Engbert–Kliegl)

Microsaccades follow Engbert and Kliegl (2003). Velocity uses their five-sample moving-average estimator, and a per-axis elliptic threshold

$$\left(\frac{\dot{x}}{\eta_x}\right)^2 + \left(\frac{\dot{y}}{\eta_y}\right)^2 > 1 \quad (3)$$

(eq. 3) flags an event when held for at least a minimum number of samples. The threshold $\eta = \lambda\sigma$ uses the **median-based** robust scale estimator $\sigma^2 = \text{med}(v^2) - \text{med}(v)^2$, clamped at zero, so a few velocity outliers cannot inflate the threshold the way a plain standard deviation would — a property the test suite checks directly (sec. 13).

6 Methods: saccade dynamics and scanpath encoding

6.1 Main sequence

For each saccade iTrace reports amplitude, direction, duration, and peak velocity. The amplitude–peak-velocity relationship — the *main sequence* — is fit to two standard parameterisations: a saturating exponential

$$V = V_{\max}(1 - e^{-A/C}) \quad (4)$$

(eq. 4; Bahill et al. (1975)) by non-linear least squares, and a power law $V = aA^b$ by linear regression in log-log space, reporting the coefficient of determination. A physiological oculomotor system gives an exponent b of roughly 0.4–0.9; deviation is a recognised marker, so the fit doubles as a sanity probe.

6.2 Scanpath encoding

Saccade sequences are encoded as direction characters following the eye-movements-as-biometrics convention: R/L/U/D for the nearest cardinal, upper-case for long saccades and lower-case for short ones relative to a length threshold. N-gram statistics over the resulting string characterise habitual scan patterns and support anomaly detection, connecting the low-level signal to higher-level behavioural analysis without leaving the pure core.

7 Methods: pupillometry

The pupil pipeline is a sequence of pure transforms over a `PupilStream`, each independently testable. It follows the conservative preprocessing vocabulary common to open pupil-analysis tools: blink handling, interpolation, baseline-correction, quality reporting, and transparent preprocessing records (Zhang and Jonides 2026; Mittner 2020). Blinks are detected as runs of NaN or sub-threshold samples and removed by linear interpolation with edge padding, so the partial-occlusion ramp on either side of a closure does not leak into the signal; a trace with no valid samples is reported as unusable rather than silently flattened to a constant. Outliers are flagged by the scaled median absolute deviation (1.4826 MAD), which is robust to the very spikes it must catch. The deblinked trace is low-pass filtered with a zero-phase Butterworth filter (with a short-trace moving-average fallback) and baseline-corrected, subtractively or divisively, against a chosen pre-event time window.

The runtime `PupilConfig` now controls the same decisions the standalone functions expose: the low-confidence pupil threshold, interpolation padding, and Butterworth cutoff/order. Reports include both cleaned-signal summaries (mean/standard deviation/min/max, phase peak and trough counts) and raw-quality summaries (valid fraction, blink fraction/count, median sample interval, peak dilation velocity, and peak constriction velocity). The live webcam pupil channel therefore remains a relative proxy, but the analysis path records how much of the trace was usable and how strongly the cleaned pupil changed over time. Brightness and contrast effects are not modelled away in this version; tools such as Open-DPSM show why dynamic visual confounds need explicit modelling when pupil size is interpreted cognitively (Cai et al. 2024).

For closed-loop and online use, a causal `PhaseDetector` classifies each streamed sample as dilation, constriction, peak, or trough using only the current and prior samples (Kronemer et al. 2024). Causality is not merely asserted but verified: feeding the detector a prefix of a signal reproduces exactly the labels it assigns to that prefix in a full run, proving no future sample influences a past label (sec. 14).

8 Methods: the capture shell

The capture shell is the sole hardware-facing module, and it is deliberately thin. Its testable heart, `iris_landmarks_to_sample`, converts one frame of normalised MediaPipe Face Mesh landmarks — supplied as a plain float array — into a binocular gaze sample, with no MediaPipe import anywhere in the call. MediaPipe Iris is cited here as a commodity-camera landmark source, not as a guarantee of calibrated eye-tracker accuracy (Vakunov et al. 2020). Because the function consumes plain floats, the entire landmark→gaze mathematics is unit-tested headless, and the same function is the bridge the forward-model loop (sec. 9) exercises.

`WebcamSource` wraps the genuinely device-bound work: opening the camera with OpenCV, running MediaPipe inference per frame, and yielding timestamped capture samples with gaze, a relative pupil/iris proxy, frame index, an FPS estimate, and quality flags. The live-frame path preserves that existing sample API while also returning the frame dimensions, a clamped pixel-space eye-region box derived from the Face Mesh eye and iris landmarks, and a JPEG data URI containing the zoomed eye crop. The proxy is explicitly labelled relative: MediaPipe Face Mesh does not expose a calibrated pupil boundary or millimetre diameter, so the live path records a camera-dependent size signal rather than a physical pupil measure. `itrace record` writes the detected gaze samples to CSV, can write the relative pupil proxy to a second CSV, and can write

the full capture-record table (`frame_index`, monotonic timestamp, gaze, pupil proxy, FPS estimate, and quality flags) via `--records-out`. `itrace camera-probe` exercises dependency and camera access without turning hardware availability into a scientific validation claim. Both hardware commands suppress noisy native OpenCV/MediaPipe stderr diagnostics by default while retaining `--backend-logs` for debugging. Constructing it without the `capture` extra installed raises a clear, actionable error naming the missing extra, rather than a bare `ModuleNotFoundError` from deep in the stack. Only the few lines that actually grab frames are excluded from coverage; the dependency-absent error path, landmark conversion, capture-sample shape, split gaze/pupil CSV writers, and full capture-record CSV writer are all tested headless.

The separate `itrace live-html` command adds a local single-user FastAPI orchestrator around the same verified Python path. The browser page receives WebSocket messages containing the eye crop, the latest capture sample, rolling `pipeline.analyze_session` summaries, event records, and plot-ready series; the browser itself only draws controls and native Canvas/SVG diagnostics. It is configured as the `web` extra and remains import-safe: `fastapi`, `uvicorn`, `OpenCV`, and `MediaPipe` are imported only inside server or capture entry points. Persistence is explicit: without `--output-dir` the app is in-memory, and with `--output-dir` it writes CSV/JSON artifacts only when export is requested.

9 Methods: 3-D forward model and the closed loop

Detectors verified against a one-dimensional position signal are only as auditable as the path that produced that signal. To exercise the geometry/landmark path end-to-end — the part a real webcam pipeline actually runs — we add a 3-D forward model (`eyemodel`) and an animated scene (`scene`) that drive the *real* estimator from synthetic anatomy with known truth.

9.1 The eyeball-to-landmark forward model

A parameterised eyeball — centre, radius, gaze yaw/pitch, and pupil and iris radii — is rotated to a true gaze direction. The iris ring and pupil disc are placed on the sphere surface at that orientation and **perspective-projected** through a pinhole camera to the normalised, MediaPipe-shaped landmark array the capture layer consumes. The anatomical eye corners are fixed: they sit on the face, not on the rotating globe, so they do not move with gaze. Because the corners are stationary while the iris translates across the visible sclera, the iris-to-corner offset encodes gaze exactly as a real face mesh would, which is what lets the same `geometry` routine that an actual recording exercises be the one under test. The projection is a forward map from a physical pose to image coordinates; it is deliberately not the inverse the estimator computes.

9.2 The animated scene

The scene (`scene`) drives a deterministic trajectory of fixation targets joined by minimum-jerk saccades — the same velocity profile the synthetic single-saccade generator uses — together with a pupil baseline carrying Gaussian dilation events and blink intervals. At each frame it emits both the per-frame 3-D truth (gaze yaw/pitch, pupil size, blink state) and the projected landmark array, so the entire recording has a ground truth that no detector could have peeked at.

9.3 The closed loop

The **closed loop** then feeds those landmarks back through the production estimator — `iris_landmarks_to_sample` reconstructs a `GazeStream`, which the standard `pipeline` analyses into fixations, saccades, and a pupil summary — and measures recovery against the 3-D truth: gaze angular error in degrees, saccade detection against the scripted saccades, and pupil correlation with the scripted baseline.

The design is deliberate on one point: the forward model (3-D sphere + perspective projection) and the estimator (the arcsine sphere approximation of sec. 4) are **independent formulations**, not two halves of one identity. Their agreement therefore tests the geometry rather than restating it — a tautology would yield exactly zero residual — while the small non-zero residual is the price of the estimator’s approximation, an honest, bounded quantity rather than a hidden one. What this construction can and cannot establish is taken up in sec. 15 and sec. 21: it certifies that the analysis code recovers truth from a geometrically faithful but *idealized* eye, and it is silent — by construction — about corneal refraction, the kappa angle, lens distortion, and every other physical effect the pinhole model omits.

10 Methods: noise model and statistical design

10.1 Idealized noise model

To ask *how much noise the recovery tolerates*, we identify the quantity that dominates corruption. Photon shot noise, JPEG compression, pixel quantisation, and MediaPipe’s finite precision all manifest downstream as one thing: **jitter in the normalised image coordinates of the landmarks**. We model it as an *idealized* additive **i.i.d. Gaussian** of standard deviation σ (normalised image-coordinate units) on every landmark before estimation. This is explicitly simplified: it ignores

spatial correlation between adjacent landmarks, temporal correlation, heteroscedasticity (error grows at eccentric gaze and under poor lighting), and systematic bias — and it is *not* derived from sensor optics. Systematic *biases* (corneal refraction, kappa angle, lens distortion) are excluded entirely; they belong to the device-validation gap (sec. 21). Appearance-based gaze work on GazeCapture, MPIIGaze, and L2CS-Net makes the same practical point from the opposite direction: real-frame accuracy depends on dataset, camera, illumination, head pose, and calibration conditions, not just on the downstream event detector (Krafka et al. 2016; Zhang et al. 2017; Abdelrahman et al. 2022). Pupil-size measurement, which a real pipeline derives from a separate segmentation we do not implement, carries its own observation noise scaled with σ as a stated modelling assumption — so the pupil channel’s noise behaviour is *assumed*, not emergent. The model is chosen to be the simplest perturbation that is falsifiable rather than the most realistic one: it is a lower bound on the kinds of corruption a real webcam imposes, and the recovery curves it produces are best read as best-case envelopes.

10.2 Statistical design

The sweep is a **sensitivity / robustness** analysis, not a hypothesis-testing power calculation. At each noise level we run n independent, seeded closed-loop recordings (sec. 9) and measure three outcomes — gaze RMS error (deg), saccade detection F1, and pupil correlation with truth. Per level we report the mean and a 95% **bootstrap percentile** confidence interval across the n trials. We use a bootstrap rather than a symmetric Student- t interval because two of the metrics are bounded ($F1 \in [0, 1]$, correlation $\in [-1, 1]$): near 1.0 a symmetric interval can run past the bound, whereas a percentile of resampled means cannot leave the range of the in-bound observations. We then locate the noise level at which each outcome crosses a usability bound by linear interpolation between grid points. The interval describes variability of the mean across the random landmark-noise realisations sampled by the n seeded trials (a Monte Carlo standard error over the noise process), not variability from re-running the deterministic pipeline. Saccade F1 scores a detected event as a true positive when its sample interval **overlaps** a ground-truth saccade interval (no separate temporal tolerance); precision and recall are computed from the same overlap match and combined as the harmonic mean, so a detector that fragments one true saccade into several events or merges several into one is penalised rather than rewarded. Every trial seed is deterministic, so the whole sweep — and therefore every figure and table number — is reproducible. Full statistical detail is in sec. 24. The sweep is falsifiable: were recovery independent of noise, gaze RMS would be flat across the grid.

The same percentile-bootstrap discipline is reused wherever the package attaches uncertainty to a point estimate — the main-sequence exponent interval of sec. 11 and the event-rate summaries of sec. 12 — for the identical reason: it assumes no symmetry and stays inside the range the data support. None of these intervals describe a population of eyes or sessions the recording was not designed to sample; they describe sampling variability of the statistic over the realisations actually observed, and they are reported as such.

For package-level validation, `itrace.validation` separates two cases that are often conflated. Synthetic-domain validation has event truth: seeded clean, webcam-jitter, head-drift, and low-light/dropout sessions provide saccade intervals, amplitudes, directions, pupil events, blinks, and quality flags, so within-domain repeat summaries and cross-domain stability gaps can report interval-overlap precision, recall, F1, amplitude error, direction error, peak-velocity error, and pupil-validity statistics. Live webcam diagnostics do not have that oracle. Following the data-quality reporting emphasis in recent eye-tracking standards work, the live HTML interface therefore reports sampling regularity, finite gaze fraction, pupil-valid fraction, path length, dispersion, and warnings, but deliberately leaves recovery F1 undefined unless a future reference-device or public-dataset path supplies truth (Jakobi et al. 2024).

11 Methods: descriptive and distribution statistics

The `itrace.stats` package turns the typed events of sec. 5 and sec. 6 into summary statistics, fitted distributions, and spread metrics. Like the rest of the core it is pure NumPy/SciPy, takes explicit seeds wherever it resamples, and operates on whatever gaze it is given — synthetic streams with known ground truth or recorded ones. The methods below *describe* a scanpath; none of them measure real-eye accuracy, which remains the device-validation gap of sec. 21.

11.1 Descriptive summaries

Given a list of fixations or saccades, `stats.descriptive` returns the count and the mean, standard deviation, median, and inter-quartile range of the relevant per-event quantity — fixation duration for fixations; amplitude, duration, direction, and peak velocity for saccades — reusing the property arrays already vectorised by `saccades.saccade_properties`. An empty list yields a count of zero and zero-filled summaries rather than raising, so a recording in which a detector finds no events of a given type flows through the pipeline without a special case. These are ordinary sample statistics over the detected events, not estimates of an underlying population the recording was not designed to sample.

11.2 Distribution fitting and model comparison

Fixation durations and saccade amplitudes are strictly positive and right-skewed, so iTrace fits three standard positive-support families by maximum likelihood: the **gamma** and **log-normal** distributions, and the **ex-Gaussian** (a Gaussian convolved with an exponential), whose right tail has long been used to model the shape of reaction-time and fixation-duration distributions (Ratcliff 1979; Luce 1986). For the two SciPy families the location parameter is **pinned to zero** (`floc=0`) so the fit respects the positive support and the scale parameter is not traded off against a spurious shift; the ex-Gaussian is fit by numerically maximising its log-likelihood over (μ, σ, τ) from method-of-moments starting values. Models are ranked by the **Akaike and Bayesian information criteria** ($AIC = 2k - 2 \ln \hat{L}$, $BIC = k \ln n - 2 \ln \hat{L}$), which penalise the extra ex-Gaussian parameter against any gain in fit, and a lower-is-better ordering is reported alongside the raw log-likelihoods. Absolute goodness of fit is summarised by a one-sample **Kolmogorov–Smirnov** statistic $D = \sup_x |F_n(x) - \hat{F}(x)|$ between the empirical CDF and each fitted CDF; we report D as a descriptive distance and treat the companion p -value with caution, since the parameters were estimated from the same sample (the classic fitted-parameter caveat). Fitting requires a minimum number of finite positive observations and raises a **ValueError** below it, mirroring the guard in `mainsequence.fit`.

11.3 Scanpath spread metrics

The spatial extent of a scanpath is summarised by four metrics computed over fixation centroids (in degrees of visual angle). **Gaze dispersion** is the root mean square distance of fixations from their centroid. The **bivariate contour ellipse area (BCEA)** — the area of the ellipse expected to contain a stated fraction (default 68%) of fixations under a bivariate-normal assumption, $BCEA = 2\pi k \sigma_x \sigma_y \sqrt{1 - \rho^2}$ with $k = -\ln(1 - P)$ — is a standard, units-of-deg² measure of fixation stability introduced for fixational eye movements (Steinman 1965) and widely used as a low-vision fixation-stability index (Castet and Crossland 2012). The **convex-hull area** of the fixation centroids gives a distribution-free footprint of explored space via `scipy.spatial.ConvexHull` (returning zero for fewer than three non-collinear points rather than raising). Both BCEA and convex-hull area are reported only as conditional descriptors: BCEA assumes approximate bivariate normality, which a structured scanpath violates, and is flagged as such.

Two **entropy** metrics capture organisation rather than extent. *Stationary position entropy* discretises fixation centroids onto a coarse spatial grid and reports the Shannon entropy of the occupancy histogram in bits — higher when gaze is spread evenly across cells, lower when it concentrates. *Direction-transition entropy* reuses the cardinal scanpath string of sec. 6: it forms the first-order transition-count matrix between successive direction symbols and reports the entropy of the transition distribution, a compact index of how predictable the scan order is, in the spirit of information-theoretic scanpath analysis (Boccignone and Ferraro 2004). Both are normalised against the maximum entropy of their respective alphabets so values are comparable across recordings of different length, and both degenerate gracefully (zero entropy) when only one cell or symbol is observed.

11.4 Bootstrap CI on the main-sequence exponent

The power-law main-sequence fit of sec. 6 reports a point estimate of the exponent b (`power_b` from `mainsequence.fit`). To attach uncertainty, `stats.scanpath_metrics.main_sequence_exponent_ci` resamples the matched amplitude / peak-velocity pairs with replacement B times (default $B = 2000$, explicit seed), refits the log-log regression on each resample, and reports a two-sided 95% **bootstrap percentile** confidence interval on b (Efron and Tibshirani 1993). We use the percentile interval here for the same reason it is used in the noise sweep (sec. 24): it makes no symmetry assumption and stays inside the range supported by the data. The interval describes sampling variability of the exponent *across the observed saccades*; it is not a claim about a population of eyes the recording did not sample. Resamples that lose enough unique positive amplitudes to make the log-log fit ill-conditioned are discarded; if no valid resample survives, the interval collapses to the point estimate rather than producing a spurious narrow bound.

12 Methods: advanced detection and scanpath comparison

The detectors of sec. 5 fix one threshold for a whole recording and reduce a scanpath to summary scalars. Two further layers relax those choices: an `itrace.detection` module that adapts its thresholds to the data and reports movement dynamics the fixed detectors discard, and two `itrace.stats` modules that compare and resolve scanpaths over time. Like the rest of the core they are pure NumPy/SciPy, take explicit seeds wherever they resample, and operate on whatever gaze they are given — synthetic streams with known ground truth (sec. 9) or recorded ones. Nothing here measures real-eye accuracy; that remains the device-validation gap of sec. 21.

12.1 Adaptive detection and movement dynamics

The fixed-threshold I-VT of sec. 5 requires the analyst to pick a velocity cut that suits the recording’s noise floor. `detection.detect_ivt_adaptive` instead follows the data-driven scheme of Nyström & Holmqvist (Nyström and Holmqvist

2010): it seeds a peak-velocity threshold, iteratively re-estimates it from the mean and standard deviation of the samples that fall *below* the current threshold (the putative fixation/noise population), and stops when the estimate converges. The converged value, not a hand-set constant, then classifies saccades, so the same call adapts to a clean high-rate trace and a noisy webcam one without re-tuning. The converged threshold is also surfaced in `SessionReport.quality` when the configurable pipeline is run with `DetectionConfig(method="adaptive_ivt")`, so a downstream report can state which threshold actually certified the detected events.

The same configurable route also exposes a conservative merge-gap repair. When `merge_gap_s` is positive, I-VT bridges only subthreshold gaps whose duration is no longer than that value and whose run is bounded on both sides by saccadic samples. This is deliberately narrower than smoothing the whole velocity trace: it repairs one known failure mode — short capture dropouts or plateaus splitting one saccade — while preserving event boundaries elsewhere and making the chosen repair window visible in the report configuration.

Fast eye movements are followed by **post-saccadic oscillations** (PSOs) — the brief wobble of the eye as it settles — which a plain I-VT either swallows into the saccade or mislabels as a tiny second saccade. `detection.detect_pso` scans the samples immediately after each detected saccade offset for a short, lower-amplitude velocity excursion in the settling window and tags it as a PSO rather than a movement, in the spirit of the glissade/PSO class that REMoDNaV adds to the saccade/fixation dichotomy (Dar et al. 2021). iTrace does not yet claim REMoDNaV’s full four-class smooth-pursuit classification — that is named as future work in sec. 21 — so the PSO tag is reported as a labelled sub-event of the saccade it follows, not as an independent event class.

Two further quantities describe the *dynamics* the fixed detectors leave on the floor. `detection.intersaccadic_intervals` returns the gaps between successive saccade onsets (seconds), the inter-saccadic-interval distribution used throughout the eye-movement literature to characterise scanning tempo and fixational dwell (Holmqvist et al. 2011); an empty or single-saccade input yields an empty array rather than raising. `detection.saccade_peak_accelerations` differentiates the velocity signal a second time and reports the peak acceleration (deg/s^2) within each saccade interval, the second-order companion to the peak-velocity main sequence of sec. 6. Both reuse the velocity estimators of sec. 5 rather than recomputing them, so a recording’s acceleration profile is consistent with the velocities its saccades were detected from.

12.2 Scanpath similarity

Comparing two scanpaths — a recording against a template, or two sessions of the same task — needs a distance, not a scalar summary. `itrace.stats.similarity` provides three complementary ones over the cardinal direction strings of sec. 6, demonstrated on recovered scanpaths in sec. 19.

The **Levenshtein scanpath distance** is the minimum number of single-character insertions, deletions, and substitutions that turn one encoded direction string into another, computed by the standard dynamic-programming edit-distance recurrence over the two strings; it captures how much one scan *order* must be rewritten to match the other and is robust to small local insertions a fixed alignment would punish. We report both the raw edit count and a length-normalised version (distance divided by the longer string’s length) so paths of unequal length are comparable, and an empty-versus-empty comparison returns distance zero.

The **n-gram cosine similarity** lifts the comparison from order to habitual sub-patterns. Each scanpath is turned into a vector of n-gram counts over its direction string — reusing `encoding.ngram_counts` — and the two vectors are compared by cosine similarity over their shared vocabulary, giving a value in $[0, 1]$ that is high when two recordings share the same characteristic short sequences regardless of where they occur. Either vector being empty (a recording with fewer than n saccades) yields a defined similarity of zero rather than a division by zero.

The **transition matrix** is the raw material both similarities are built to summarise: the first-order count (or row-normalised probability) of moving from each direction symbol to each next one, the same construction that underlies the direction-transition entropy of sec. 11 and the generative, information-theoretic view of scanpaths (Boccignone and Ferraro 2004). Exposing the matrix directly lets an analyst inspect *which* transitions drive a similarity or entropy value rather than trusting the scalar, and the row-normalised form is guaranteed to be a proper stochastic matrix or an all-zero row for an unobserved symbol.

12.3 Windowed event-rate time series

A single recording is rarely stationary: scanning tempo, fixation stability, and blink rate drift over a task. `itrace.stats.timeseries` slices a recording into fixed-width time windows (with an optional overlap) and reports, per window, the fixation rate, the saccade rate, and the blink rate in events per second, each built from the typed events the core already detected and the blink intervals of the pupil channel. The result is an aligned set of arrays — window-centre times and the three rate series — suitable for plotting a session’s evolution or for feeding the spatial-stability metrics of sec. 11 (gaze dispersion, BCEA, used as a low-vision fixation-stability index (Castet and Crossland 2012)) a window at a time. Windows that contain no samples

report a rate of zero rather than `nan`, and a window narrower than the sampling interval is rejected with an actionable error, so a mis-specified window size fails loudly instead of producing an empty or misleading series.

13 Results: ground-truth recovery

The first question for any detector is whether it recovers events it could not have peeked at. Every trace in this section is synthetic: a generator builds the signal *and* returns the parameters used to build it, so the recovered numbers are checked against truth held out by construction. None of these results speak to real-eye accuracy, which remains the device-validation gap of sec. 21.

On a synthetic fixation→saccade→fixation trace at 250 Hz, the I-VT detector (Salvucci and Goldberg 2000) recovers exactly one saccade bracketed by two fixations; the recovered amplitude matches the constructed 10–12° to within 5% and peak velocity to within 10% across the test grid. The min-jerk generator sweeps amplitude and direction, and the recovered direction tracks the constructed angle around the full circle (the gaze convention of 0° right, +90° up), so a sign flip or axis swap on either path would surface immediately. A purely-noisy fixation trace yields zero saccades at the default threshold — the anti-criterion confirming the detector does not hallucinate events from noise.

The Engbert–Kliegl detector (Engbert and Kliegl 2003) recovers an embedded 0.5° microsaccade from fixational jitter while ignoring the surrounding noise; a probe confirms the threshold uses the median-based robust estimator (it returns a value below half the plain standard deviation in the presence of velocity outliers), the property checked directly in sec. 5. Recovered microsaccade amplitude and peak velocity fall on a tight, low-amplitude main-sequence cluster, distinct from the large-saccade branch — the expected separation of the two scales rather than a fitted claim. The main-sequence fit recovers the saturating $V_{\max}$ and C of a synthetic relationship (Bahill et al. 1975) to within 10% and returns a power-law exponent inside the physiological 0.4–0.9 range (fig. 1).




Figure 1: Main-sequence recovery from a synthetic multi-amplitude recording. Panel A plots detected saccades by amplitude, peak velocity, and direction, with the fitted saturating curve and recovered parameters. Panel B repeats the same detections on log-log axes with the fitted power law, making the exponent and goodness-of-fit explicit.

14 Results: pupillometry

A sinusoidal pupil signal with an injected NaN blink is correctly deblinked (no residual NaNs), and the interpolated-plus-smoothed trace coincides with the underlying sine away from the blink window — the blink is bridged, not fabricated. A fully-invalid trace is reported as unusable rather than silently flattened, so a recording the camera never resolved cannot masquerade as a flat pupil. These are recovery checks against a known generator, not statements about pupil-size accuracy on a real eye (sec. 21).

The causal phase detector labels peaks within a two-sample window of the analytic sine maxima, and the online-equals-offline-prefix test confirms it consults no future sample: replaying the signal one sample at a time reproduces the offline labels exactly up to each prefix, so the causality property of sec. 7 holds empirically rather than by assertion. The detected peak

and trough counts match the number of sine cycles in the trace, the coarsest sanity bound on the detector. Saccade direction distributions over the same synthetic sessions are summarised as polar histograms (fig. 2).




Figure 2: Saccade direction distribution over the synthetic multi-saccade recording. The polar convention is printed in the figure (0° = right, $+90^\circ$ = up), bars encode counts over 16 angular bins, and the orange resultant vector summarises directional balance.

15 Results: closed-loop recovery

Running the full loop — 3-D eyeball $\rightarrow$ perspective projection $\rightarrow$ landmark extraction $\rightarrow$ estimation $\rightarrow$ recovery — on the default animated scene (four fixations joined by three saccades, two pupil dilation events, one blink, at 120 Hz) recovers gaze with an RMS residual of **0.16°** against the 3-D truth (maximum 0.23°) for gaze within $\pm 15^\circ$, and recovers the three scripted saccades exactly (fig. 3). The residual is small but strictly non-zero, confirming the forward model and estimator are independent (sec. 9).

This is an *internal-consistency verification*: it proves the estimator correctly inverts the forward model and would expose sign flips, axis swaps, unit errors, and coordinate-frame mismatches. It is **not** device validation. The 0.16° is a *lower bound* on real-world error — any idealization the two paths share (pinhole projection, no corneal refraction, a spherical eye, zero kappa) contributes exactly zero residual by construction (sec. 21). The recovered pupil ($r = 1.0$) is a near-identity consistency check (the projected pupil/iris ratio is a monotone function of pupil radius), confirming the deblink/pipeline preserve the signal — not pupil-size accuracy.

The synthetic generator supplies all three target signals jointly — gaze direction, saccade intervals, and pupil diameter — animated as floating 3-D eyeball orbs (fig. 4).

16 Results: noise-sensitivity analysis

Sweeping the landmark-noise standard deviation σ from 0 to 0.016 (normalised image units), 25 seeded trials per level, gives a clear and differentiated degradation of the three signals (fig. 5, tbl. 1).




Figure 3: Closed-loop recovery on the default synthetic scene. Panel A shows the projected binocular landmark geometry and the two iris-centre paths. Panel B overlays true and recovered yaw/pitch, with scripted and detected saccade intervals shaded. Panel C overlays true and recovered pupil dynamics after normalisation, with the blink interval shaded.



Figure 4: Synthetic truth and recovered signals rendered together. The 3-D panel shows the true eyeball orientation, gaze rays, iris, and pupil state at one representative frame; the trace panels repeat the recovered yaw, pitch, and pupil signals with true values, saccades, and the blink interval visible for audit.

Gaze RMS error rises monotonically and approximately linearly with σ , crossing the 2° usability bound at $\sigma \approx 0.005$. **Saccade** detection is the most fragile emergent signal: F1 falls below 0.8 already at $\sigma \approx 0.0014$ — a direct consequence of detection resting on the velocity, the time-derivative of position, which amplifies high-frequency noise. Both orderings *emerge* from propagating landmark noise through the real estimator.

The **pupil** result carries a sharp caveat: its robustness (correlation above 0.9 until $\sigma \approx 0.005$) is **robust by construction** under the assumed `pupil_noise_scale`, not a measured property, and would move if that parameter changed. It is reported as a conditional illustration, not a finding; the genuine emergent result is the gaze-vs-saccade ordering.

The most defensible practical takeaway comes from translating σ to pixels. At a 640 px width, the saccade breakdown at $\sigma \approx 0.0014$ is ≈ 0.9 px, whereas gaze tolerates roughly three pixels in this idealised sweep. That pixel-scale comparison is not a universal detector specification: real webcam landmark error is camera-, pose-, lighting-, compression-, and algorithm-dependent. It does show why webcam saccade *timing* is theoretically marginal before any real-world bias, temporal correlation, or tracking dropout is added, while coarse fixation and gaze direction have more headroom.




Figure 5: Recovery vs idealised webcam observation noise. Each panel reports the Monte-Carlo mean and 95% bootstrap confidence band for one recovered signal, with the manuscript usability bound drawn as a horizontal line and the interpolated bound crossing annotated in both normalised σ and approximate pixels at 640 px width. Colour, marker, and line style all encode the signal so the figure remains legible in greyscale; the pupil panel is robust only under the stated `pupil_noise_scale` assumption.

The same numbers, with per-cell uncertainty, are the canonical statistics table (tbl. 1); it and the figure are generated from one sweep so they never drift. (Snapshot at $n = 25$; regenerate with `scripts/generate_power_figure.py`.)

Table 1: Recovery vs landmark noise σ (mean $\pm$ 95% half-CI, $n = 25$).

σ (norm)	σ (px@640)	gaze RMS (deg)	saccade F1	pupil r
0.0000	0.00	0.163 $\pm$ 0.000	1.000 $\pm$ 0.000	1.000 $\pm$ 0.000
0.0010	0.64	0.452 $\pm$ 0.006	0.949 $\pm$ 0.029	0.996 $\pm$ 0.000
0.0020	1.28	0.842 $\pm$ 0.012	0.612 $\pm$ 0.028	0.982 $\pm$ 0.001
0.0040	2.56	1.665 $\pm$ 0.026	0.220 $\pm$ 0.007	0.934 $\pm$ 0.003
0.0080	5.12	3.418 $\pm$ 0.046	0.274 $\pm$ 0.017	0.800 $\pm$ 0.008
0.0160	10.24	6.798 $\pm$ 0.102	0.560 $\pm$ 0.043	0.546 $\pm$ 0.018

17 Results: cross-implementation agreement and gates

On a shared synthetic trace the core I-VT and microsaccade detectors agree with an independent reference implementation (`reference_impl`, raw finite-difference velocity and an explicit-loop microsaccade detector) on event count and onset timing, despite using different velocity formulations — a guard against a bug shared by a single implementation. The two paths share only their input and the published algorithm definitions, so agreement is evidence the implementation matches the algorithm, not that two copies of one mistake agree.

The complete toolchain is green and reproducible: 429 no-mocks tests pass with 91.18% statement-and-branch coverage, **ruff** and **ruff format** report no issues, and **mypy --strict** finds no type errors across the source tree. Every synthetic generator and every Monte-Carlo trial is explicitly seeded with `numpy.random.default_rng`, so each figure and table number is reproducible byte-for-byte from a clean checkout — re-running the figure pipeline twice produces identical PNGs, and the noise sweep of sec. 16 reproduces its table cells exactly. No test mocks a computation: detectors run on real arrays with known ground truth, file I/O uses real temporary files, and figures are asserted to write non-empty PNGs. The full module map, dependency contract, and test strategy are in sec. 25.

The validation-domain suite extends that reproducibility contract from a single trace to stress domains. **itrace synthetic-validation** writes JSON for repeated clean, webcam-jitter, head-drift, and low-light/dropout sessions, summarising within-domain recovery and across-domain stability. The live HTML page calls the same Python route when the user runs synthetic validation from the browser; the browser only renders the returned statistics and does not reimplement event detection or recovery scoring.

18 Results: figure gallery

The analysis core is paired with a deterministic visualization layer (**itrace.viz**, matplotlib on the headless Agg backend) that renders the standard eye-movement figures directly from a **SessionReport**. Each figure is produced from a synthetic session whose ground truth is known by construction, so the gallery is reproducible byte-for-byte and never depends on a captured recording. Every panel draws from one shared Wong (2011) colour-blind-safe palette and a common house style — a readable font floor, white background, and clean spines — applied uniformly across the gallery and the manuscript figure scripts so all published figures match; in the multi-series comparison figures, signals are separated by marker and line style as well as colour, so the figures stay legible in greyscale. The figures below are regenerated by **itrace figures --out-dir output/figures --animations; scripts/generate_figures.py** also refreshes the static gallery.

The gallery spans the three signal families iTrace recovers. For **gaze and events**, a velocity trace (**output/figures/velocity_trace.png**) overlays the 2-D speed signal with the I-VT threshold and shades each detected saccade, while a spatial scanpath (**output/figures/scanpath.png**) draws fixations sized by dwell time and joined by saccade arrows in screen coordinates. For **saccade dynamics**, an amplitude histogram (**output/figures/amplitude_histogram.png**) carries its maximum-likelihood distribution overlay, and a main-sequence panel (**output/figures/main_sequence_diagnostics.png**) plots amplitude against peak velocity on log-log axes with the fitted power law and its residuals. For **pupillometry**, a pupil trace (**output/figures/pupil_trace.png**) shows the raw signal, the blink-interpolated and smoothed overlay, and the causally-detected dilation peaks and troughs.

A multi-panel session dashboard (**output/figures/session_dashboard.png**) composes the velocity trace, scanpath, amplitude distribution, main sequence, saccade-direction polar histogram, and a textual summary into a single at-a-glance figure. The gallery now also includes gaze density, fixation heatmap, AOI dwell, event raster, pupil PSD, saccade-rate, microsaccade, and deterministic synthetic replay outputs. Because every panel is a pure function of a **SessionReport**, the same plots apply unchanged to a real recording analysed through the pipeline.

19 Results: scanpath comparison and temporal dynamics

The descriptive scanpath metrics of sec. 11 and the spatial/timeline visualisations of sec. 18 turn a single recording into a comparable summary of *where* and *when* gaze moved. As everywhere in this section, the inputs are synthetic sessions with known ground truth: the numbers below are internal-consistency checks on the metrics and detectors, not measurements of a real eye (sec. 21).

19.1 Self-similarity is a tautology, used as a control

Comparing a synthetic session against an exact copy of itself returns a similarity of 1.0, and comparing it against a deterministically time-shifted copy returns the same value once the paths are aligned. This is true *by construction* — identical fixation centroids and saccade sequences cannot differ — so it is reported not as a finding but as the degenerate anchor every similarity measure must hit. Its value is as a negative control: a measure that failed to return 1.0 on a session against itself would be broken, and perturbing one fixation by a known offset drops the similarity below 1.0 in the expected direction, confirming the measure responds to real differences rather than always saturating.

19.2 Adaptive and fixed thresholds recover the same scripted events

The scripted multi-saccade sessions are built so that fixed-threshold I-VT (Salvucci and Goldberg 2000) and a data-adaptive threshold derived from the velocity distribution (Nyström and Holmqvist 2010) should classify the same intervals. They do: on the synthetic grid both recover the same saccade count with onsets agreeing to within a sample, and the adaptive threshold

settles inside the band of fixed thresholds that bracket the clean velocity gap. Where the signal carries injected noise the two diverge first on the smallest events, the same fragility that the saccade-F1 collapse of sec. 16 makes quantitative. The point is a cross-check between two independent detector designs on a shared, known-truth signal — the kind of agreement an event-classification reference implementation is meant to provide (Dar et al. 2021) — not a claim that either threshold is correct for real webcam data.

19.3 Windowed event rates expose temporal structure

Sliding a fixed-width window across a session and counting events per window yields a saccade-rate and fixation-rate time series whose envelope tracks the scripted event schedule: rate peaks line up with the constructed saccade bursts and fall to zero across the long fixations, and integrating the windowed counts recovers the total event counts of the whole-session report. The microsaccade rate over a fixational-jitter session stays near zero except around the single injected microsaccade, the temporal counterpart to the amplitude separation seen in sec. 13. These windowed rates are descriptive summaries of the detected event stream; on a synthetic session they recover the schedule we wrote in, and on a real recording they would inherit exactly the detector fragilities catalogued in sec. 16, with no additional accuracy implied.

Taken together, the spatial spread metrics (dispersion, convex-hull area, BCEA, and the two scanpath entropies of sec. 11) and the timeline views give a consistent, reproducible picture of the synthetic sessions: structured scanpaths read as low direction-transition entropy and a compact hull, diffuse ones as the opposite, and the temporal rates localise each event in time. Every number is a property of the detected events on a known-truth signal — a verification that the metrics and detectors behave as specified, kept strictly separate from any claim about real-eye behaviour.

20 Discussion

20.1 What the decoupling buys

By implementing the canonical algorithms in a hardware-free core and verifying them against constructed ground truth, iTrace makes the scientific layer continuously testable and reproducible. A regression in the saccade detector, the microsaccade threshold, or the blink interpolation is caught by a headless test run, not discovered later on real data. The optional capture shell can then evolve — new landmark backends, browser capture, deep-learning gaze — without putting the verified core at risk, because the shell only produces the gaze and pupil streams the core already analyses. The local HTML orchestrator is an example of this discipline: the browser receives eye crops and plot-ready summaries, but all capture, estimation, event detection, statistics, and export remain in the same Python package path tested elsewhere.

20.2 What the expanded analysis layer adds

The same decoupling extends past detection into description. On top of the estimators sits a quantitative layer (`itrace.stats`) and a visualization layer (`itrace.viz`), both pure-core in spirit: the statistics modules import only NumPy/SciPy, and matplotlib is confined to the optional `figures` extra so the core stays headless. The statistics layer turns the typed events of sec. 5 and sec. 6 into descriptive summaries of fixation and saccade quantities, maximum-likelihood fits of positive-support distributions (gamma, log-normal, ex-Gaussian) ranked by AIC/BIC with a Kolmogorov–Smirnov distance, scanpath spread and organisation metrics (gaze dispersion, convex-hull area, BCEA, and spatial / direction-transition entropy), and a bootstrap-percentile confidence interval that attaches sampling uncertainty to the main-sequence exponent (sec. 11). The visualization layer renders the standard velocity, scanpath, distribution, main-sequence, microsaccade, and pupil figures — and a composite session dashboard — deterministically and byte-for-byte from a `SessionReport` (sec. 18).

The honest framing matters here too. Every one of these methods *describes* a scanpath it is handed; none of them measures real-eye accuracy or repairs the device-validation gap of sec. 21. A tighter bootstrap interval, a lower KS distance, or a smaller BCEA is a property of the *detected events*, not evidence that those events match a real eye. What the layer genuinely adds is reach and reproducibility: the same descriptive vocabulary other eye-movement packages expose, computed by tested, seed-pinned functions over the verified core, applicable unchanged to a synthetic stream with known ground truth or a real recording pushed through the pipeline. The most useful comparison point is therefore not “is iTrace more accurate than WebGazer, L2CS-Net, REMoDNaV, pymovements, PupEyes, or pypillometry?” but “does iTrace make the exact algorithmic contract clear enough that those systems can be benchmarked against it on shared fixtures or public datasets?” (Papoutsaki et al. 2016; Abdelrahman et al. 2022; Dar et al. 2021; Krakowczyk et al. 2023; Zhang and Jonides 2026; Mittner 2020).

20.3 Algorithmic verification, not device validation

A precise claim matters. What sec. 13 establishes is *algorithmic verification*: each estimator recovers the parameters of signals produced by a known forward model. This is not *device validation* — agreement with an independent reference measurement of the same physical quantity on real eyes. We have not performed the latter. A passing synthetic suite proves the estimation

mathematics is correct and regression-protected; by itself it is compatible with poor real-world performance. The hardware capture path is not exercised by the suite, and it is exactly where the dominant error sources live: head-pose and calibration confounds in gaze, the pupillary light reflex in pupillometry, and saccade undersampling at ~ 30 Hz. Until a head-to-head against a reference tracker or a public dataset exists, iTrace should be cited as a verified reference implementation of the algorithms, not a validated tracker.

20.4 The closed loop bounds, but does not establish, real error

The 3-D closed loop (sec. 15) exercises the landmark $\rightarrow$ gaze projection, but its guarantee is bounded by what the forward model and estimator *share*. Both assume an ideal pinhole projection with no lens distortion, no corneal refraction (the apparent pupil differs from the physical pupil by $\sim 0.5\text{--}1$ mm), a perfectly spherical eyeball, and a zero kappa angle ($\sim 5^\circ$ in real eyes). Any error these shared idealizations introduce is identically zero in the loop by construction, and the synthetic landmarks carry none of MediaPipe’s real noise, bias, or eccentric-gaze foreshortening. The 0.16° residual is thus a *lower bound* on real-world error, and the $r = 1.0$ pupil result is a consistency check, not accuracy. The quantitative limits this implies are collected in sec. 21.

21 Limitations and accuracy bounds

The conveniences of the pure core do not change the physics of consumer hardware, and several limits should temper interpretation. The headline limitation is the unclosed **device-validation gap**: iTrace’s correctness claims rest on constructed ground truth and an internal-consistency loop, never on agreement with a reference measurement of real eyes. Everything below follows from that.

Gaze and saccades. Real-frame webcam gaze accuracy depends strongly on the dataset, camera, lighting, pose, calibration, and model class; large public benchmarks report errors in physical screen units or degrees under their own protocols rather than a universal webcam constant (Krafka et al. 2016; Zhang et al. 2017; Abdelrahman et al. 2022). iTrace’s noise-sensitivity result sharpens the method-level risk: sec. 16 shows saccade *timing* breaks down at a landmark perturbation of ≈ 0.9 px, so velocity-thresholded webcam saccade timing is theoretically marginal before real-world lens, illumination, head-pose, landmark-bias, and calibration effects are added.

Pupil. Webcam pupil diameter now has dedicated learned-model and dataset work (Shah et al. 2025), while downstream packages emphasise transparent preprocessing, visualization, and statistical modelling once pupil samples exist (Zhang and Jonides 2026; Mittner 2020). iTrace does not implement an absolute pupil segmentation or millimetre calibration path, so its live pupil channel remains a relative proxy. The package surfaces this rather than implying IR-grade precision, and the pupil noise-sensitivity result is conditional on a modelling assumption (sec. 10), not measured.

The expanded analysis layer describes, it does not validate. The descriptive statistics, distribution fits, scanpath spread and entropy metrics, and the bootstrap confidence interval of sec. 11 all operate on the events a detector returns. They quantify the *shape* of a scanpath precisely, but a crisp gamma fit, a narrow exponent interval, or a small BCEA says nothing about whether the underlying gaze tracked a real eye — it is description downstream of the same unvalidated estimates, not independent corroboration of them.

The live HTML app is a diagnostic interface, not a validation study. A large zoomed eye crop, stable FPS counter, and plausible rolling saccade/pupil plots prove that the webcam, MediaPipe, and Python analysis path are wired together. They do not prove gaze accuracy, pupil-diameter accuracy, or event timing agreement with an independent reference device.

Idealizations the validation cannot see. The closed loop and the noise model share assumptions — pinhole optics, no refraction, spherical eye, zero kappa, iid landmark noise — and are blind to any error those introduce (sec. 20). Spatial/temporal landmark-noise correlation, heteroscedasticity, and systematic bias are unmodelled.

21.1 Future work

The honest path to closing the device-validation gap is empirical: a head-to-head against a reference eye-tracker or a public webcam dataset (MPIIGaze, GazeCapture) on real frames, reporting error in degrees and pupil correlation. Beyond that, an absolute-diameter pupil model trained on PupilSense/EyeDentify data (Shah et al. 2025), smooth-pursuit and post-saccadic-oscillation classification to match REMoDNv’s four-class output (Dar et al. 2021), screen-space calibration, and binocular disparity each slot in behind the existing typed interfaces without disturbing the verified core.

22 Conclusion

iTrace shows that auditable webcam eye-movement analysis follows from a single architectural commitment: keep the verifiable algorithms pure and verify them against ground truth, and keep the fragile hardware in a thin, optional shell. The

science is dependable *because of*, not in spite of, that separation. On the verified core sits an analysis surface — descriptive event statistics, distribution fitting and comparison, scanpath spread and entropy metrics, a bootstrap interval on the main-sequence exponent, and a deterministic figure gallery — that extends the package’s descriptive reach without weakening any correctness claim, because each layer is itself tested and seed-pinned and describes the scanpath it is given rather than asserting real-eye accuracy. The contribution is honestly scoped — a verified, type-safe, reproducible reference implementation of the estimation algorithms and their descriptive statistics, with a 3-D-forward-model internal-consistency check and an idealized noise-sensitivity analysis whose most defensible result is that webcam saccade timing is theoretically marginal under sub-pixel-to-few-pixel landmark perturbations. What remains is the empirical step: real frames, a reference device, and the upper bound on error those would finally establish.

23 References

- Abdelrahman, Ahmed A., Thorsten Hempel, Aly Khalifa, and Ayoub Al-Hamadi. 2022. “L2CS-Net: Fine-Grained Gaze Estimation in Unconstrained Environments.” *arXiv Preprint arXiv:2203.03339*, ahead of print. <https://doi.org/10.48550/arXiv.2203.03339>.
- Bahill, A Terry, Michael R Clark, and Lawrence Stark. 1975. “The Main Sequence, a Tool for Studying Human Eye Movements.” *Mathematical Biosciences* 24 (3-4): 191–204. [https://doi.org/10.1016/0025-5564\(75\)90075-9](https://doi.org/10.1016/0025-5564(75)90075-9).
- Boccignone, Giuseppe, and Mario Ferraro. 2004. “Modelling Gaze Shift as a Constrained Random Walk.” *Physica A: Statistical Mechanics and Its Applications* 331 (1-2): 207–18. <https://doi.org/10.1016/j.physa.2003.09.011>.
- Cai, Yuqing, Christoph Strauch, Stefan Van der Stigchel, and Marnix Naber. 2024. “Open-DPSM: An Open-Source Toolkit for Modeling Pupil Size Changes to Dynamic Visual Inputs.” *Behavior Research Methods* 56: 5605–21. <https://doi.org/10.3758/s13428-023-02292-1>.
- Castet, Eric, and Michael Crossland. 2012. “Quantifying Eye Stability During a Fixation Task: A Review of Definitions and Methods.” *Seeing and Perceiving* 25 (5): 449–69. <https://doi.org/10.1163/187847611X620955>.
- Dar, Asim H, Adina S Wagner, and Michael Hanke. 2021. “REMoDNaV: Robust Eye-Movement Classification for Dynamic Stimulation.” *Behavior Research Methods* 53 (1): 399–414. <https://doi.org/10.3758/s13428-020-01428-x>.
- Efron, Bradley, and Robert J Tibshirani. 1993. *An Introduction to the Bootstrap*. Chapman & Hall/CRC. <https://doi.org/10.1201/9780429246593>.
- Engbert, Ralf, and Reinhold Kliegl. 2003. “Microsaccades Uncover the Orientation of Covert Attention.” *Vision Research* 43 (9): 1035–45. [https://doi.org/10.1016/S0042-6989\(03\)00084-1](https://doi.org/10.1016/S0042-6989(03)00084-1).
- Holmqvist, Kenneth, Marcus Nyström, Richard Andersson, Richard Dewhurst, Halszka Jarodzka, and Joost van de Weijer. 2011. *Eye Tracking: A Comprehensive Guide to Methods and Measures*. Oxford University Press.
- Jakobi, Deborah N., Daniel G. Krakowczyk, and Lena A. Jäger. 2024. “Reporting Eye-Tracking Data Quality: Towards a New Standard.” *Proceedings of the 2024 Symposium on Eye Tracking Research & Applications*, 47:1–3. <https://doi.org/10.1145/3649902.3655658>.
- Kassner, Moritz, William Patera, and Andreas Bulling. 2014. “Pupil: An Open Source Platform for Pervasive Eye Tracking and Mobile Gaze-Based Interaction.” *Proceedings of the 2014 ACM International Joint Conference on Pervasive and Ubiquitous Computing: Adjunct Publication*, 1151–60. <https://doi.org/10.1145/2638728.2641695>.
- Krafka, Kyle, Aditya Khosla, Petr Kellnhofer, et al. 2016. “Eye Tracking for Everyone.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2176–84. https://openaccess.thecvf.com/content_cvpr_2016/html/Krafka_a_Eye_Tracking_for_CVPR_2016_paper.html.
- Krakowczyk, Daniel G et al. 2023. “Pymovements: A Python Package for Eye Movement Data Processing.” *arXiv Preprint arXiv:2304.09859*, ahead of print. <https://doi.org/10.48550/arXiv.2304.09859>.
- Kronemer, Sharif I et al. 2024. “Cross-Species Real-Time Detection of Trends in Pupil Size Fluctuation.” *Behavior Research Methods*, ahead of print. <https://doi.org/10.3758/s13428-024-02374-8>.

- Luce, R. Duncan. 1986. *Response Times: Their Role in Inferring Elementary Mental Organization*. Oxford University Press.
- Mittner, Matthias. 2020. “Pypillometry: A Python Package for Pupillometric Analyses.” *Journal of Open Source Software* 5 (51): 2348. <https://doi.org/10.21105/joss.02348>.
- Nyström, Marcus, and Kenneth Holmqvist. 2010. “An Adaptive Algorithm for Fixation, Saccade, and Glissade Detection in Eyetracking Data.” *Behavior Research Methods* 42 (1): 188–204. <https://doi.org/10.3758/BRM.42.1.188>.
- Papoutsaki, Alexandra et al. 2016. “WebGazer: Scalable Webcam Eye Tracking Using User Interactions.” *Proceedings of IJCAI*, 3839–45.
- Ratcliff, Roger. 1979. “Group Reaction Time Distributions and an Analysis of Distribution Statistics.” *Psychological Bulletin* 86 (3): 446–61. <https://doi.org/10.1037/0033-2909.86.3.446>.
- Salvucci, Dario D, and Joseph H Goldberg. 2000. “Identifying Fixations and Saccades in Eye-Tracking Protocols.” *Proceedings of the 2000 Symposium on Eye Tracking Research & Applications*, 71–78. <https://doi.org/10.1145/355017.355028>.
- Shah, Vijul, Ko Watanabe, Brian B. Moser, and Andreas Dengel. 2025. “PupilSense: A Novel Application for Webcam-Based Pupil Diameter Estimation.” *arXiv Preprint arXiv:2407.11204*, ahead of print. <https://doi.org/10.48550/arXiv.2407.11204>.
- Steinman, R. M. 1965. “Effect of Target Size, Luminance, and Color on Monocular Fixation.” *Journal of the Optical Society of America* 55 (9): 1158–65. <https://doi.org/10.1364/JOSA.55.001158>.
- Vakunov, Andrey, Dmitry Lagun, Chuo-Ling Chang, and Matthias Grundmann. 2020. *MediaPipe Iris: Real-Time Iris Tracking and Depth Estimation*. Google Research Blog. <https://research.google/blog/mediapipe-iris-real-time-iris-tracking-depth-estimation/>.
- Zhang, Han, and John Jonides. 2026. “PupEyes: An Interactive Python Library for Eye Movement Data Processing.” *Behavior Research Methods* 58 (1): 29. <https://doi.org/10.3758/s13428-025-02830-z>.
- Zhang, Xucong, Yusuke Sugano, Mario Fritz, and Andreas Bulling. 2017. “MPIIGaze: Real-World Dataset and Deep Appearance-Based Gaze Estimation.” *arXiv Preprint arXiv:1711.09017*, ahead of print. <https://doi.org/10.48550/arXiv.1711.09017>.

24 Supplement S1: statistical methods

This supplement gives the full statistical detail behind sec. 16 so the analysis is reproducible and auditable.

24.1 Estimator and interval

For each (noise level, metric) we collect n independent trial values $\{x_i\}$ and report the sample mean $\bar{x}$ with a two-sided 95% **bootstrap percentile** confidence interval: we draw $B = 2000$ resamples of size n with replacement, take the mean of each, and read the 2.5th and 97.5th percentiles of the resampled means.

We deliberately avoid a symmetric Student- t interval $\bar{x} \pm t_{0.975, n-1} s / \sqrt{n}$ here because two of the three metrics are **bounded** — saccade $F1 \in [0, 1]$ and pupil correlation $\in [-1, 1]$ — and at low noise both sit near 1.0, exactly where a symmetric interval can extend past the bound and report a nonsensical value above 1. A percentile of resampled means can never leave the range of the in-bound observations, so the interval is always valid. For the unbounded gaze RMS the two methods agree closely.

24.2 What the interval means (and the determinism caveat)

The pipeline is fully deterministic: re-running it reproduces every number byte-for-byte (sec. 17). The confidence interval therefore does **not** describe “what happens if I rerun the code.” It describes the variability of the mean **across the random landmark-noise realisations sampled by the n seeded trials** — a Monte Carlo standard error of the mean over the noise process. With $n = 25$ fixed seeds it quantifies how much the level-mean would move under a different finite draw of noise realisations.

24.3 Trial independence and determinism

Each trial uses a distinct deterministic seed (`base_seed + level_index*1000 + trial`), and the bootstrap resampler is itself fixed-seeded, so trials are independent draws of the landmark-noise process while the whole sweep — including every CI — is byte-for-byte reproducible.

24.4 Saccade detection scoring

Saccade F1 uses **interval-overlap** matching against the ground-truth saccade intervals: a true saccade is recalled if any detected interval overlaps it, and a detected interval is a true positive if it overlaps any true saccade. There is no separate temporal-tolerance parameter; precision, recall, and $F_1 = 2PR/(P + R)$ follow directly. This makes the score insensitive to small onset/offset shifts while still penalising spurious or missed events.

24.5 Threshold crossings

A usability threshold (gaze RMS = 2° , saccade F1 = 0.8, pupil $r = 0.9$) is located by linear interpolation between the two grid points that bracket the crossing; with $n = 25$ these crossings are approximate and should be read to ~ 2 significant figures. The crossing in physical units uses `sigma_to_pixels( $\sigma$ , image_width)` (sec. 10).

24.6 One source of truth

The figure (fig. 5) and the table (tbl. 1) are both rendered from a single `NoiseSweepResult` via `power.summary_records / format_summary_markdown`, so prose, plot, and table cannot drift apart — verified by a cell-for-cell match in the documentation audit.

25 Supplement S2: software architecture

25.1 Module map

Module	Layer	Responsibility
<code>config</code>	core	frozen configs for detection, pupil, capture, figures
<code>types</code>	core	dataclasses + units (<code>GazeStream</code> , <code>PupilStream</code> , events)
<code>geometry</code>	core	pix $\leftrightarrow$ deg, iris $\rightarrow$ angle, interocular normalisation
<code>calibration</code>	core	affine gaze calibration and recording-quality summaries
<code>velocity</code>	core	Savitzky–Golay / gradient velocity
<code>saccades</code>	core	I-VT, I-DT, Engbert–Kliegl microsaccades
<code>detection</code>	core	adaptive I-VT, PSO candidates, ISI, acceleration
<code>mainsequence</code>	core	saturating + power-law fit
<code>pupil</code>	core	deblink, MAD, low-pass, baseline
<code>pupilphase</code>	core	causal phase detector
<code>encoding</code>	core	direction-character + n-gram scanpath
<code>eyemodel</code>	core	3-D eyeball + pinhole projection
<code>scene</code>	core	animated trajectory + closed loop
<code>synthetic</code>	core	seeded gaze/pupil sessions with truth records
<code>power</code>	core	noise-sensitivity sweep + statistics
<code>stats</code>	core	descriptive, distribution, similarity, event, bootstrap metrics
<code>pipeline / io</code>	core	end-to-end analysis + CSV
<code>capture</code>	shell	OpenCV + MediaPipe (lazy)
<code>live</code>	shell	local FastAPI/WebSocket HTML orchestrator (lazy)
<code>dashboard</code>	shell	Streamlit + Plotly (lazy)
<code>viz</code>	shell	matplotlib figure gallery (lazy via figures extra)
<code>cli</code>	shell	typer command-line interface

25.2 Dependency contract

The shell may import the core; the core may never import the shell. Optional dependencies (`opencv-python`, `mediapipe`, `fastapi`, `uvicorn`, `httpx2`, `streamlit`, `plotly`, `matplotlib`) are imported lazily inside functions or the optional `itrace.viz` subpackage, never by `import itrace`, and absence raises an actionable error rather than breaking import. `import itrace` and the full test suite run with none of the hardware/dashboard/web dependencies installed; the `all` extra installs capture, dashboard, figure, and web backends together for local smoke testing.

25.3 Test strategy and gates

Tests are **no-mocks**: each detector is verified by synthesising a signal whose ground truth is known by construction and asserting recovery within a stated tolerance, plus an independent reference implementation for cross-checking (sec. 17). Promotion gates: `pytest` green at $\geq 90\%$ statement+branch coverage, `ruff check` and `ruff format --check` clean, and `mypy --strict` clean. Figures and tables are regenerated by thin orchestrator scripts under `scripts/` that contain no analysis logic of their own.